Evaluating experiments

Andy J. Wills

Spurious correlations

http://www.tylervigen.com/spurious-correlations

Depression and memory

- Depression is associated with over-general memory.

- Depression causes memory problems?
- Memory problems cause depression?
- Both causal directions?
- Neither causal direction (e.g. both caused by childhood trauma).

- It is not possible to distinguish between these accounts on the basis of correlational data.

Longitudinal data does not solve this problem

- Use of night lights in infancy is correlated with myopia in later life (true).
- Seems causal? Causes must precede effects. The later myopia cannot cause the earlier use of night lights. So, night lights must be causing myopia?
- Ban night lights? (genuinely recommended on basis on these data).

Third factor explanations are still possible in longitudinal research

- A third factor causes both the presence of night lights and myopia.
- Developing myopia in later life has a genetic component. If your parents are myopic, this increases the chance you will become myopic.
- Myopic adults, on average, favour higher levels of illumination. This drives their decision to use night lights in their baby's room.
- The parents' myopia causes both the presence of infant night lights and later myopia.
- Ban night lights? Clearly, this would be ineffective.

Correlation does not imply causation

- Correlational research is fundamentally limited.
- It is extremely unlikely that any two variables are completely unrelated.
- Many correlations in psychology are very small e.g.
  - Extroversion explaining 2% of the variation in some other variable.
  - 2% is detectably different from no correlation
  - but not meaningful (everything likely to be related to some degree).

Determining causation through the Experimental Method

- Simplest form
  - Take two groups of people
  - Do different things to those two groups.
  - Measure something

- Independent variable - Intended difference in what we do to the two groups
- Dependent variable - The thing we measure

Example: Testing a treatment for depression

- Group 1 - 6 weeks of the new therapy
- Group 2 - Nothing.
- Take measure of depression at end (e.g. Beck Depression Inventory).
- Group 1 are less depressed than Group 2

- This has the potential to show that the therapy *causes* a reduction in depression.
- ...but there are other explanations.

Pre-existing differences

- Group 1 - 6 weeks of the new therapy
- Group 2 - Nothing.

- What if Group 1 were happier to start with?

- Approaches to this problem
  - Detection
  - Prevention

Detection

- Take pre-treatment measures
- e.g. Measure BDI of both groups before (and after) treatment period.

	Pre	Post
Therapy	25	5
Control	25	25

Prevention

- Construct groups such that we eliminate pre-existing differences.
- Matching - Take BDI measures for everyone. Allocate people to groups in such a way that the average BDI for the two groups is identical (or at least, minimised).
- Randomisation - Allocate people to groups randomly.
- Matching versus Randomisation - pros and cons.

Our therapy experiment

- Use large, randomised groups.
- Take pre-treatment measures
- Treatment caused the reduction in depression?

	Pre	Post
Therapy	25	5
Control	25	25

Attrition

- Attrition - participants dropping out before the end of the study
- If attrition rates vary between conditions, you may have a major problem.

Example

- Pre-treatment BDI scores

						Mean
Therapy	6	8	12	15	30	14.2
Control	6	8	12	15	30	14.2

- The most-depressed 20% drop out of therapy (perhaps because the therapy is quite demanding).
- There are no drop-outs in the control condition (there's not much to drop out from).
- Both therapy and control are inert (no effect) - post-treatment BDI equals pre-treatment BDI.

Example

- Pre-test BDI scores

						Mean
Therapy	6	8	12	15	30	14.2
Control	6	8	12	15	30	14.2

- Post-test BDI scores

						Mean
Therapy	6	8	12	15		10.25
Control	6	8	12	15	30	14.2

- A therapy we know to be ineffective appears to have worked, due to non-random attrition.

Placebo effect

- Classic example
  - Someone has a headache
  - Give them a pill with no active ingredient
  - Tell them it's a headache tablet
  - Their headache symptoms reduce
- Lesson - In order to assess drug effectiveness you need to test drug vs. placebo, NOT drug vs. nothing.

Placebo effect in psychological therapy

- Perhaps the therapy is inert?
- The treatment group are happier because they have the expectation that what they are receiving will work.
- Problem - a placebo pill is known to be inert; what is the equivalent in therapy?
- There is no agreement - there's someone willing to endorse the effectiveness of almost any therapy.

Placebo effect in psychological therapy

- Solution - set out to show that your new therapy works better than an existing treatment (or, as well as existing treatment, if yours is better in some practical way e.g. cheaper).
- Problem - this is seldom done.

Experimenter Effects - Data analysis - Example

- Diary entries as a measure of happiness.
- Participants write about their feelings
- Experimenter rates for level of happiness.
- If experimenter knows which condition the participant is in, this may bias their assessment of happiness.

Experimenter Effects - Data analysis

- Machine-recorded measures (e.g. reaction time) immune?
- No! - Data analysis typically involves many decisions, all open to bias.
- If the experimenter knows which condition the participants are in, this could bias their decisions.

Blind testing

- Single-blind testing - participant does not know which condition they are in.
  - e.g. Drug vs. placebo. Participants do not know which condition they are in.
- Double-blind testing - single-blind testing plus the experimenters do not know which condition is which until after they have completed their analysis.

Pre-registration

"The first principle is that you must not fool yourself, and you are the easiest person to fool" - Richard Feynman.

- Record your hypothesis, method, and analysis plan, before you analyse the data.

Difference versus no difference designs

- The preferred hypothesis is that people differ in the speed with which they react to auditory and visual alarm signals.
- The alternative theory against which this is compared is that there is no difference (nil hypothesis).
- Problem - Experimental control is never perfect.
- Thus - the nil hypothesis is almost certainly wrong, and detectably so if you test enough people.
- Thus - the result of the study is known before you run it.
- Thus - There was no point in running it.

Better alternatives 1

- Directional hypotheses
  - The preferred theory is that auditory is faster.
  - The alternative theory against which this is compared is that there is no difference (nil hypothesis).
  - If you find visual faster, you have disproved your theory.
  - So, whatever the result, there was a point to running this experiment (because the theory was falsifiable).

Better alternatives 2

- Strong inference
  - One well-established theory predicts that auditory is faster.
  - Another well-established theory predicts that visual is faster.
  - Whatever you find in this study, you've gained information (except in the unlikely case where the nil hypothesis was true).

Evaluating an experiment

Mueller, P. A., & Oppenheimer, D. M. (2014). The pen is mightier than the keyboard: Advantages of longhand over laptop note taking. *Psychological science, 25*, 1159-1168.

1. Find the full text of this paper on Google Scholar
2. Read from the title up to, but not including, the "Study 2" subheading.
3. Evaluate how good Study 1 is, using the *checklist* to help remind you of what we've covered today.
4. Agree on a score, be ready to report your score, and to answer some questions.

Further reading/ watching

The notes for this lecture cover a number of additional relevant topics.